

---

# Getting Started with Node

# Tools

---

## Tools you need:

- NodeJS – basically a Google V8 Javascript engine
  - available from nodejs.org, easy install
- Npm – the Node Package Manager
  - Available from <https://www.npmjs.com/>
  - Helps you install and manage 3<sup>rd</sup> party modules

## Optional tools you don't need but might like:

- Nvm – the Node Version Manager
  - Available via npm install -g nvm, or by download from nodejs.org
- Bower - “front-end package manager” – you shouldn't need
- Gulp – build system. Kinda like ant, but you write Javascript tasks in a “gulpfile” and then run.
- Grunt – ditto, just goes about its business differently

## What about an IDE?

- Yes, you can use Eclipse, get the nodeclipse plugin
- Or there are several good JS editors and run node on the CL

## First example:

---

Hello World:

```
console.log("Hello World!");  
$ node hello.js
```

Um yeah...maybe a little more interesting:

```
/* From theprojectspot.com */  
var http = require('http');
```

Load module, like "import" (kinda)

Create the http server, and on each request execute the callback

```
http.createServer(function (req, res) {  
    res.writeHead(200, {'Content-Type': 'text/plain'});  
    res.end('Hello World\n');  
}).listen(8080);
```

What port (and optionally hostname) we listen on

```
console.log('Server started');
```

*Well that was less code & setup  
than a servlet → Stay tuned...*

# Functional Programming

---

A review of SER221:

- In functional programming, you encapsulate behaviors not state
- Unlike procedural programming, there are no side-effects
- It has seen a bit of resurgence lately
  - Languages include Haskell, LISP, Clojure, F#
  - Javascript and Scala walk the line between functional and OO
  - Java8 even added lambda expressions to move this way

And this matters for web development because?

- The functional approach means we layer *functional* application semantics over our request/response synchronous HTTP
  - In other words, we model a single request/response as a pipelined process of discrete stateless computational tasks
- Compared to OO (like Java), we avoid the overhead of object construction/deconstruction and impedance mismatch

# Asynchronous Processing

---

In the conceptual overview, we talked about the asynchronous single-threaded execution model

- Single-threaded
- Non-blocking
- Use of callbacks

## Pros and Cons of the Asynchronous Model

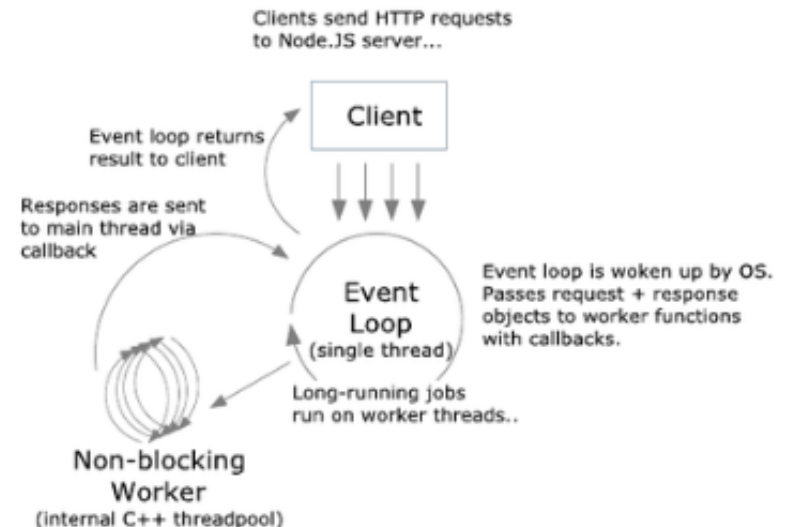
### Pros:

1. No “synchronized” blocks
2. Event loop is efficient
3. No context switch

### Cons:

1. Awkward coding model at first
2. Threads natural extension of process
3. Not natural to OO

*It is more natural to think in terms of a functional, task-oriented paradigm, where each distinct computation is executed off the “queue”*




# Functional & Asynchronous Programming

---

*...these go together like...*

The event queue is important

- It is that constant line of work (think treadmill) that the single thread keeps chewing through



How do you manipulate it?

- The natural flow is to provide a callback
  - From a functional design perspective, you have to make sure work is broken into discrete chunks
    - Your function, on completion, puts more work on the queue
  - From an asynchronous perspective, you have to realize you are not sequential, but you are not concurrent either



There are other ways of adding work to the queue:

1. Make a blocking I/O call that ends up on a worker thread
2. Use an event listener for a built-in event
3. Emit your own event and write custom listeners
4. Use `process.nextTick` as a scheduling mechanism
5. Use explicit timers to schedule work



# Explicit Timers

---

3 types of timers:

1. 1-shot timeout: `setTimeout`

2. Periodic: `setInterval`

3. Immediate: `setImmediate` – after I/O callbacks but before “regular” work

```
// simpleblock.js
// Example from Ch. 2 of Mixu's Node Book, http://book.mixu.net/node/ch2.html
var myTimer = setTimeout(function() {
    console.log('Timeout at ' + new Date().toString());
}, 500);
myTimer.unref(); // if timers are all that is on the event queue, exit
// Change the function between setTimeout, setInterval, and setImmediate
// start time
var sTime = new Date();
console.log('Started app processing loop at ' + sTime.toString());

// delay block
var i = 0;
while (new Date().getTime() < sTime.getTime() + 2500) { i++ }
console.log('Exiting processing loop at ' + new Date().toString() + ' after
    ' + i + ' iterations');
```

# Process.nextTick()

---

This gets your work on the event queue before blocking I/O

- But you could starve I/O if you are not careful
- It is slightly different from `setImmediate` in that it will invoke its callback before even going back to the event queue.
- Technically what it does is insert itself as the 1<sup>st</sup> thing done after the current event loop completes by placing its work on the *next tick queue*
- Most common usage is to ensure truly async behavior
- **Example: `simplenexttick.js`**

```
console.log('start');
process.nextTick(() => {
  console.log('nextTick callback');
});
console.log('scheduled');
function hello() {
  console.log('hello world');
}
hello();
process.nextTick(hello);
```

## **Output:**

```
start
scheduled
hello world
nextTick callback
hello world
```



# NextTick example

---

Here is a nexttick brain teaser – what is the output?

```
// nexttick.js - what will be the order of the output?
setTimeout(function timeout() {
  console.log('TIMEOUT FIRED');
}, 0)
process.nextTick(function () {
  process.nextTick(function () {
    console.log(1);
    process.nextTick(function () { console.log(2); });
    process.nextTick(function () { console.log(3); });
  });
  process.nextTick(function () {
    console.log(4);
    process.nextTick(function () { console.log(5); });
    process.nextTick(function () { console.log(6); });
  });
});
```

# Event listeners

---

```
// modified from http://www.hacksparrow.com/node-js-eventemitter-tutorial.html
// see simpleEventListener.js
var EventEmitter = require('events').EventEmitter;
var radium = new EventEmitter();

function cb(ray) {
    console.log(ray);
}

// on and addListener are the same thing
radium.on('radiation', cb);
radium.on('foo', function(ray) {
    console.log("Boo hoo I will never get called");
});
radium.once('radiation', function(ray) {
    console.log("JUST ONCE " + ray);
});
setTimeout(function() {
    radium.removeListener('radiation', cb);
}, 5000);
setInterval(function() {
    radium.emit('radiation', 'GAMMA');
}, 500);
```

## Set callbacks

- One for radiation
- One that is never called
- Another for radiation that will only get called once

Make a 1-shot timeout to remove our listener

Make an interval time out that generates events

## More on events

---

Often we will listen on events that come from built-in node modules, like in module `http`.

In the previous example we created our own event generator, and *emitted* events

- At the top we new'd up an `EventEmitter`

To do this in general, you modify the prototype of your object so it can emit events directly:


```
function myObj() { /* do your constructor work here */ }  
myObj.prototype.__proto__ =  
    require('events').EventEmitter.prototype;  
// Now you can just do this:  
myObj.emit('myEvent');
```

See `emitter_listener.js`

# Blocking I/O and the event queue

Node realizes that the most common thing a program will do that blocks is I/O

- Traditionally, we read in data from a file or write data out to a file and simply wait for the operation to complete
- If the file I/O is a long-running operation, we either tell the user to wait, or we get cute and put it in a thread
- Recognizing the danger of blocking I/O, and that it is about the only thing it knows would block ahead of time, node knows it needs to farm it off to a worker thread and provide a callback
- The 5<sup>th</sup> way you can manipulate the event queue
- Also note that because it is special, this is why we have things like `process.nextTick` that pre-empt I/O calls on the queue



# [A]Synchronous File I/O

---

You can do File I/O synchronously, which basically looks like every other programming language's simple file I/O libraries.

- Some distinctions in asynch file operations
  - Exceptions are caught by node and returned as a 1<sup>st</sup> param err object
  - Asynch functions require extra param for a callback
  - Synchronous file I/O starts right away before any subsequent events are processed by node; whereas asynch file I/O puts the next behavior on the event queue for deferred execution.

Functions: `(var fs = require('fs'))`

- Opening and Closing files:
  - `fs.open[Sync](path, flags [,mode] {,callback})` – standard file open call. Path is either relative or absolute, flags determines, read, write, append, etc. Mode is the file access mode (default rw) and callback is for async only.
  - `fs.close[Sync](fd {,callback})` – closes the file described by file descriptor fd. Closing flushes any buffers

# [A]Synchronous File I/O

---

## More file I/O Operations:

- `writeFile[Synch](path, data, [options], callback)`
  - data can be String or a Buffer object
  - writes the entire file in one swoop (See *file\_write.js*)
- `readFile[Synch]`
  - reads entire file (See *file\_read.js*)

Note when we read/write an entire file we may have a long-running operation. To keep with the principle of short discrete computational tasks, we can read/write bytes at a time:

- `write[read][Synch](fd, data, offset, length, position, callback)` – reads bytes either into data (Buffer) or writes bytes to it (String or Buffer) (See *file\_read[write]\_[a]sync*)
- You can also get file metadata. (See *file\_stats.js*, *file\_readdir.js*, *file\_folders.js*)

# [A]Synchronous File I/O

---

## Streaming I/O

- A stream is an I/O abstraction that
  - Is defined by its I/O properties not by its physical implementation
  - Works well in a *pipeline* way of thinking about your I/O
- Node provides the Stream module to working with streaming I/O
  - Emits events when data is written or read
  - We saw this already with the low-level http module – requests and responses are instances of streams (example: see `http_server_post.js`)

## A 4<sup>th</sup> way to read/write from files is to treat them as Streams

- `fs.createReadStream(path [, options])`
- `fs.createWriteStream(path [, options])`
  - Options is just an object that has configuration information as before, like encoding and access mode, etc.)
  - Your code should listen for the 'data', 'end' and 'err' events in its event listeners ("on" and "end" methods)
  - See `file_write_stream.js` and `file_read_stream.js`

# Intro to NodeJS summary

---

We are just getting started!

- Embrace the functional programming paradigm
- Embrace the asynchronous model – it takes a while for it to feel natural, so you have to practice!

Node is a modular, functional, asynchronous platform

- We will be using modules like http and events often
  - We may even write a few of our own – very easy!
- Functional – think of your tasks, not your objects
- Asynchronous – lots of callbacks, you get used to “chaining” your callbacks to get that semi-sequential feel
  - But realize other functions may try to manipulate the queue
- 5 ways to manipulate the asynchronous event queue
- Keep your tasks small, discrete, and well-defined!
  - Node wants to run the treadmill at the highest speed!